

02562 Rendering - Introduction

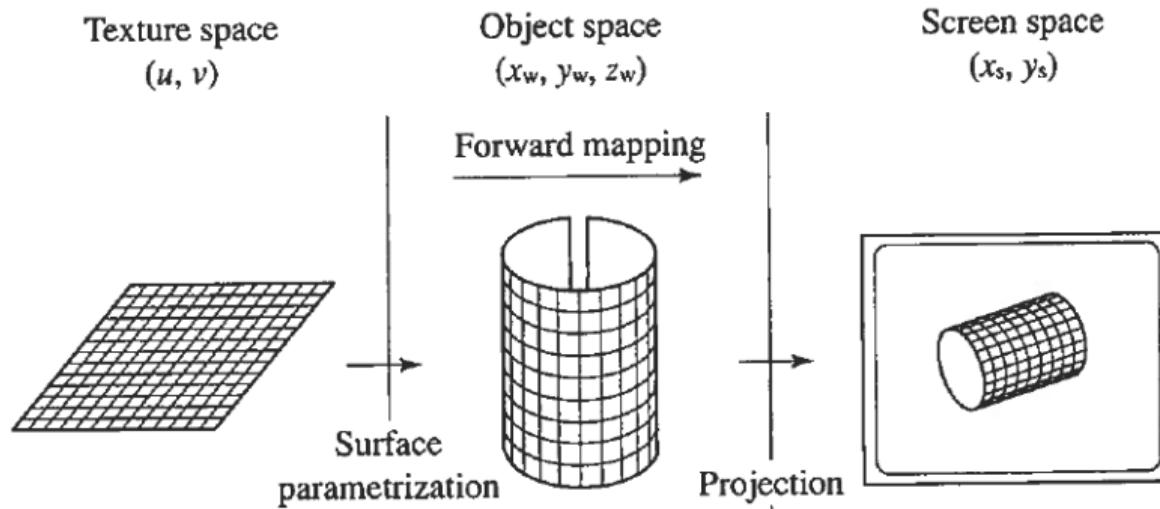
Texture Mapping

Jeppe Revall Frisvad

September 2023

Texture in computer graphics

- ▶ A texture is a 1- to 3-dimensional array of vector elements.
- ▶ Each element in the array is called a texel (texture element).
- ▶ Texture mapping is to map the texels onto an object in 3D space and then project them back into 2D image space.



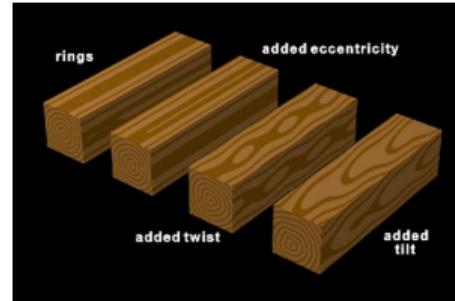
- ▶ The texels are used to modulate the object properties.

Obtaining textures

- ▶ Textures are usually loaded from image files.
- ▶ There are numerous online resources
(e.g. <https://texturehaven.com/>, <https://www.textures.com/>)
- ▶ If you want to repeat the texture across a surface, you need a *tileable* texture.
- ▶ Tileable versus non-tileable textures:

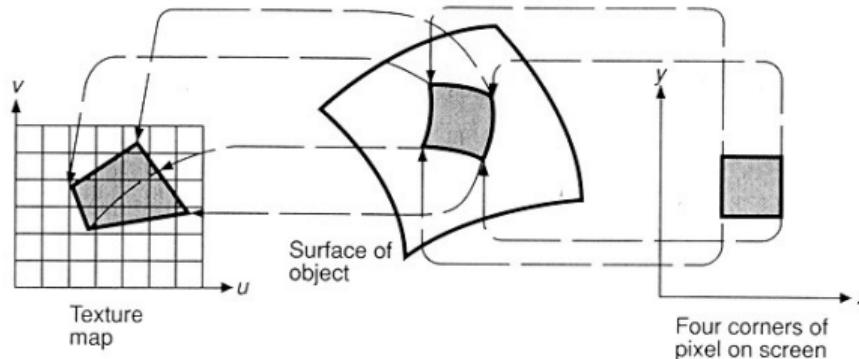


- ▶ It is also possible to generate textures procedurally.
- ▶ Example of procedural solid texturing:



Texture coordinates

- ▶ We use texture coordinates (u, v) to map a 2D texture onto the surface of a 3D object.



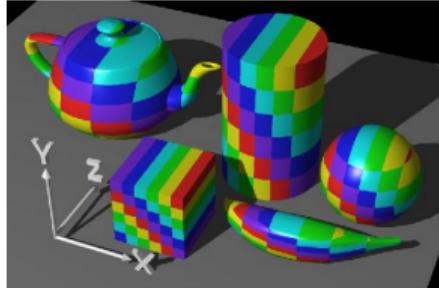
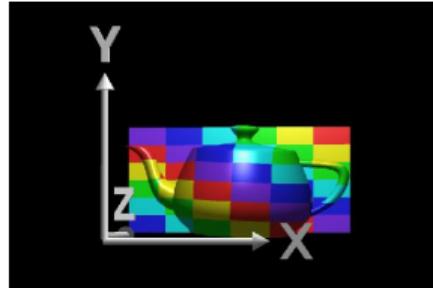
- ▶ This is a parametrization of the surface, where the parameters (u, v) map to a point on the surface in 3D space.
- ▶ The parametrization of a plane is $\mathbf{x} = \mathbf{x}_0 + u\vec{b}_1 + v\vec{b}_2$, where $\vec{n} = \vec{b}_1 \times \vec{b}_2$ is the normal of the plane while \vec{b}_1 and \vec{b}_2 are orthogonal vectors (tangent and binormal) in the plane.
- ▶ In ray tracing, we find a ray-surface intersection point \mathbf{x} . To find the texture coordinates (u, v) of an intersection point \mathbf{x} , we need an inverse mapping.

Mapping examples

- ▶ Planar map:

$$(u, v) = s(x, y)$$

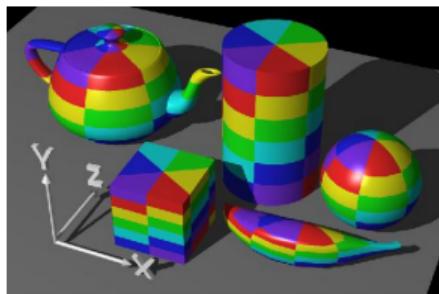
s is a scale.



- ▶ Cylindrical map:

$$(u, v) = \left(\frac{\theta}{2\pi}, h \right)$$

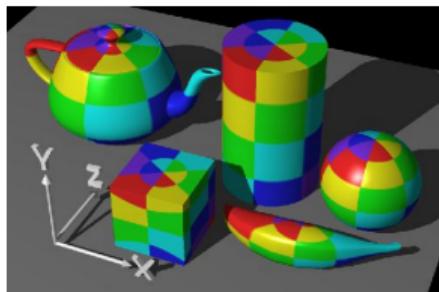
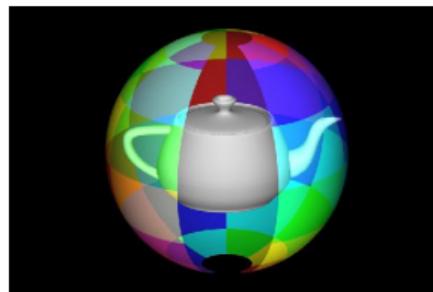
$$\theta \in [0, 2\pi], h \in [0, 1].$$



- ▶ Spherical map:

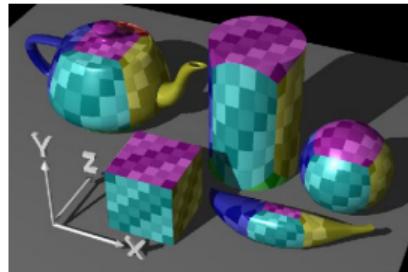
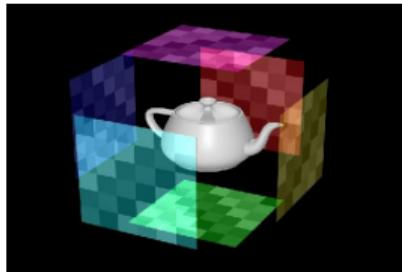
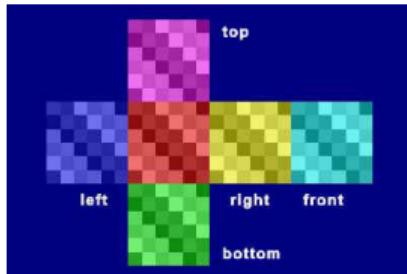
$$(u, v) = \left(\frac{\theta}{\pi}, \frac{\varphi}{2\pi} \right)$$

$$\theta \in [0, \pi], \varphi \in [0, 2\pi].$$



Cube map and environment mapping

- ▶ Cube map:

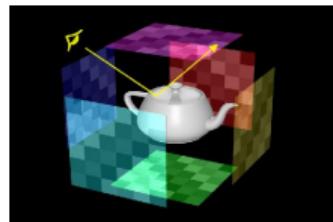
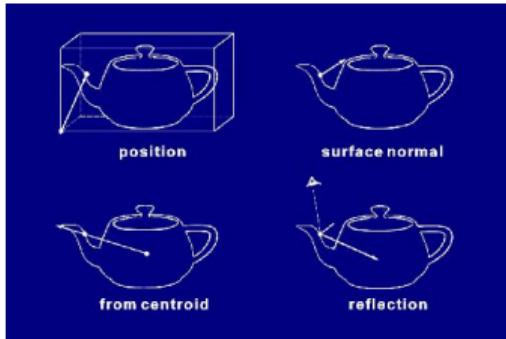


- ▶ Choosing a map entity (examples):

- ▶ position relative to the object's bounding box,
- ▶ surface normal at the point being rendered,
- ▶ vector from the centroid through the point,
- ▶ reflection vector at the current point.

- ▶ Environment mapping

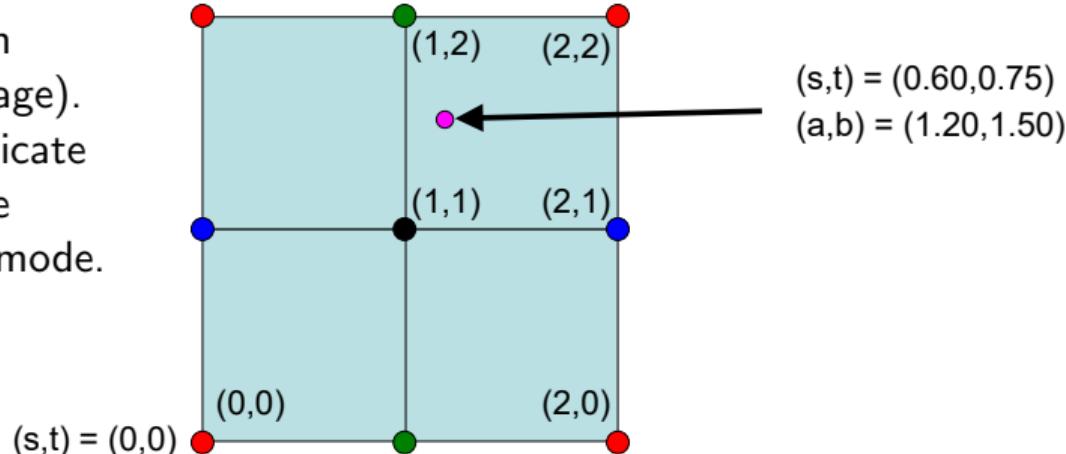
- ▶ Using ray direction for look-ups when no object was hit.
- ▶ Using the reflection entity for reflective objects.



Texture look-up

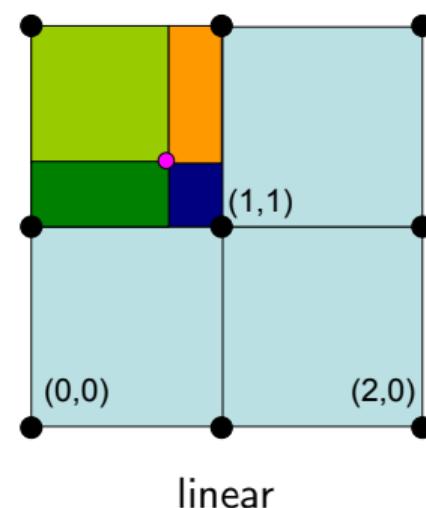
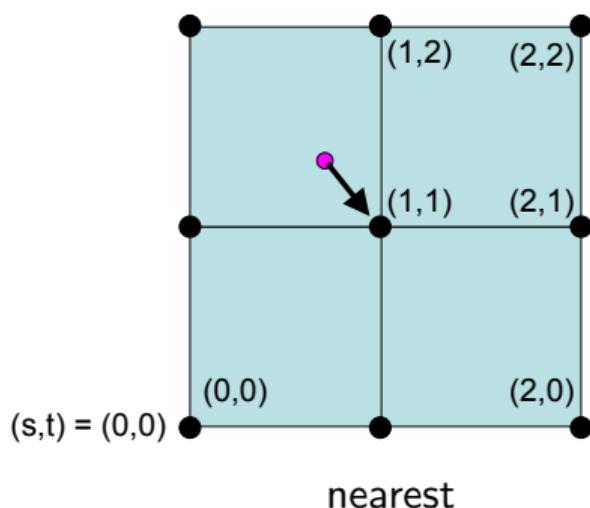
- ▶ From texture coordinates (u, v) to texture image index i .
- ▶ Two new sets of coordinates:
 - ▶ (s, t) texture space $[0, 1] \times [0, 1]$.
 - ▶ (a, b) image space $[0, W] \times [0, H]$, where $W \times H$ is the image resolution.
- ▶ In texture repeat mode: $s = u - \text{floor}(u)$, $t = v - \text{floor}(v)$.
 $a = s * W$, $b = t * H$.

- ▶ Illustration
(2×2 image).
Colors indicate
texel reuse
in repeat mode.



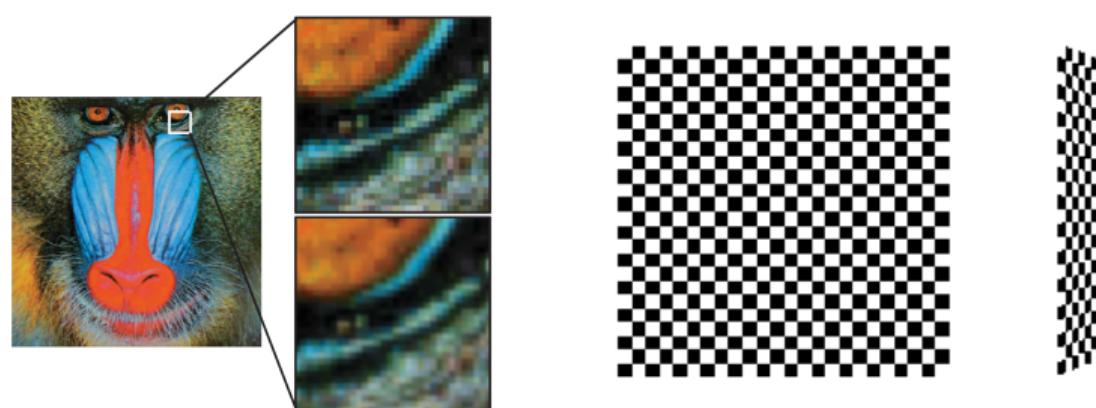
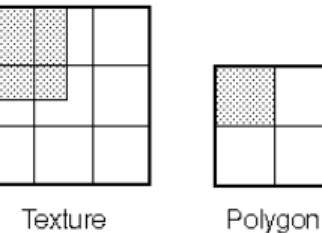
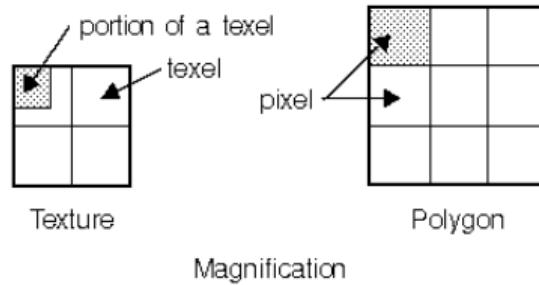
Texture filtering

- ▶ Find index i of nearest texel (nearest neighbor filtering):
$$U = (\text{int})(a + 0.5) \bmod W, \quad V = (\text{int})(b + 0.5) \bmod H, \quad i = U + V * W.$$
- ▶ Find four nearest texels, do weighted average (bilinear filtering):
$$U = (\text{int})a, \quad V = (\text{int})b, \quad c_1 = a - U, \quad c_2 = b - V,$$
where c_1 and c_2 are the weights in the bilinear interpolation.



Texture aliasing

- ▶ Individual texels rarely correspond to individual pixels of the final image.
 - ▶ Magnification is when a texel covers multiple pixels.
 - ▶ Minification is when a pixel covers multiple texels.

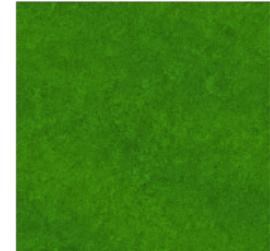


Drawing a texture using WebGPU

- ▶ A minimalistic fragment shader:

```
@group(0) @binding(0) var my_sampler: sampler;
@group(0) @binding(1) var my_texture: texture_2d<f32>

@fragment
fn main_fs(@location(0) coords: vec2f) -> @location(0) vec4f
{
    let uv = coords*0.5 + 0.5;
    return textureSample(my_texture, my_sampler, uv);
}
```



- ▶ Setup on the JavaScript side:

```
const texture = await load_texture(device, "data/grass.jpg");
const bindGroup = device.createBindGroup({
    layout: pipeline.getBindGroupLayout(0),
    entries: [
        { binding: 0, resource: texture.sampler },
        { binding: 1, resource: texture.createView() }
    ],
});
function animate() { render(device, context, pipeline, bindGroup); }
animate();
```

- ▶ The render function creates a command encoder to set and submit a render pass.
- ▶ Call with a different bindGroup using a different sampler to change texture mode.

Loading a texture for WebGPU

- ▶ Loading is asynchronous and uses HTML elements:

```
async function load_texture(device, filename)
{
    const img = document.createElement("img");
    img.src = filename;
    await img.decode();

    const imageCanvas = document.createElement('canvas');
    imageCanvas.width = img.width;
    imageCanvas.height = img.height;
    const imageCanvasContext = imageCanvas.getContext('2d');
    imageCanvasContext.drawImage(img, 0, 0, imageCanvas.width, imageCanvas.height);
    const imageData = imageCanvasContext.getImageData(0, 0, imageCanvas.width, imageCanvas.height);
```

- ▶ Copying the image data to a Uint8Array (reversing the vertical axis):

```
let textureData = new Uint8Array(img.width*img.height*4);
for(let i = 0; i < img.height; ++i)
    for(let j = 0; j < img.width; ++j)
        for(let k = 0; k < 4; ++k)
            textureData[(i*img.width + j)*4 + k] = imageData.data[((img.height - i - 1)*img.width + j)*4 + k];
```

- ▶ Streaming the image data to the GPU:

```
const texture = device.createTexture({
    size: [img.width, img.height, 1],
    format: "rgba8unorm",
    usage: GPUTextureUsage.COPY_DST | GPUTextureUsage.TEXTURE_BINDING
});
device.queue.writeTexture({ texture: texture }, textureData,
    { offset: 0, bytesPerRow: img.width*4, rowsPerImage: img.height, }, [img.width, img.height, 1]);
```

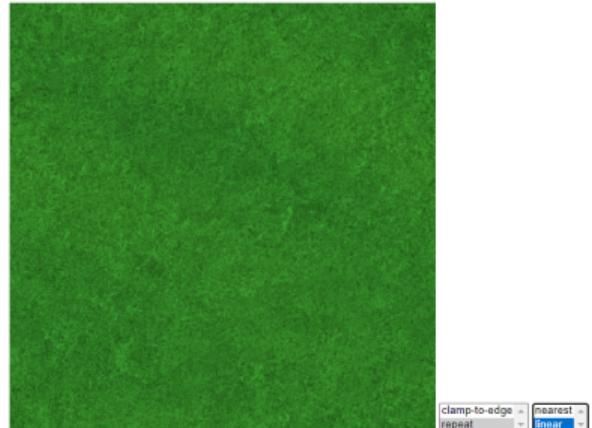
Texture addressing and filtering modes

- ▶ Conclude the `load_texture` function by creating a sampler:

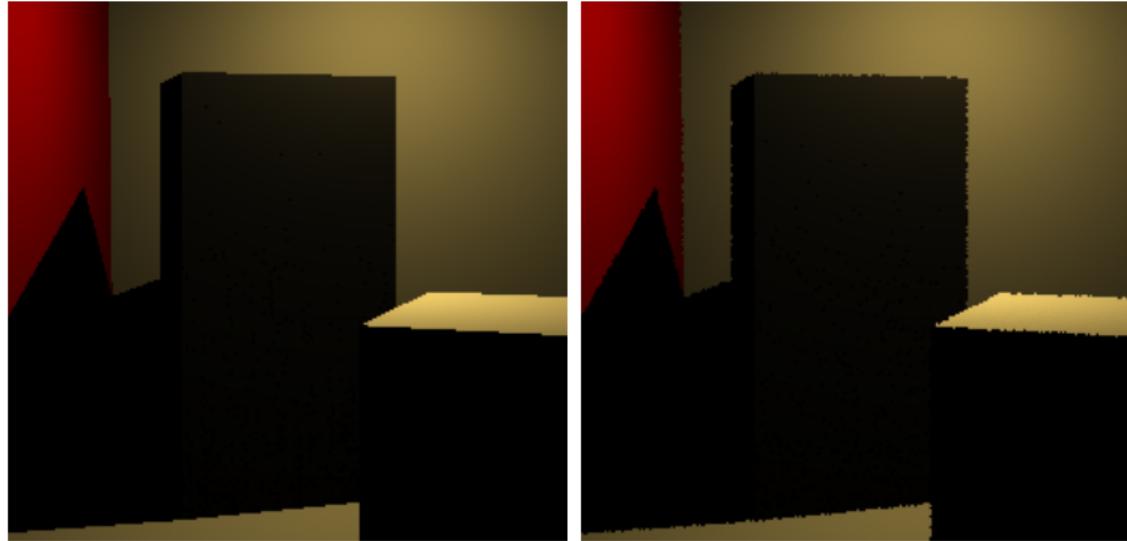
```
texture.sampler = device.createSampler({  
    addressModeU: "clamp-to-edge",  
    addressModeV: "clamp-to-edge",  
    minFilter: "nearest",  
    magFilter: "nearest",  
});  
return texture;  
}
```



- ▶ Another option for an address mode is `repeat`.
- ▶ Another option for a filter mode is `linear`.
- ▶ Create an array of samplers if you would like to switch between different modes.

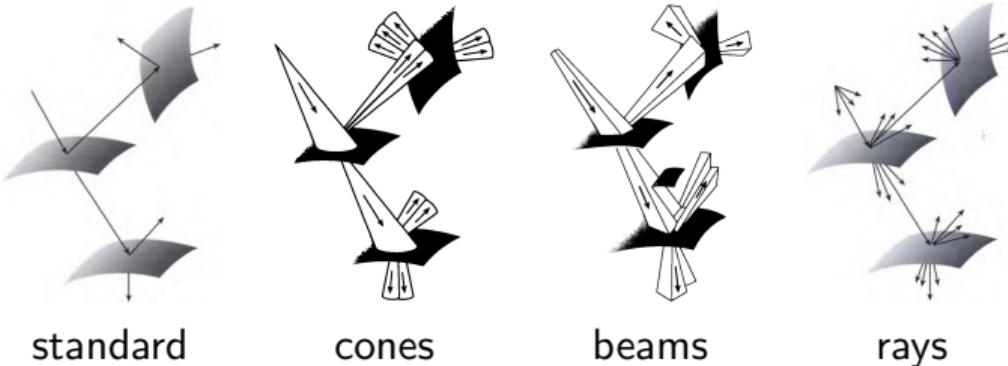


Aliasing artifacts



- ▶ Staircase artifacts ('jaggies') are obvious for hard shadows, but occur for all edges.
- ▶ Regularity artifacts appear when we always trace our rays through the pixel center.
- ▶ Trying with rays through random pixel positions leads to noise. It's a trade-off.
- ▶ How do we find out if a pixel covers multiple surfaces or texels?

Pixel coverage?



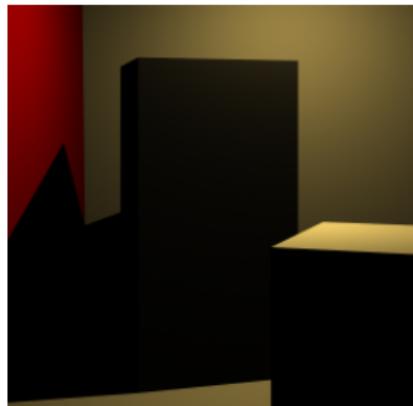
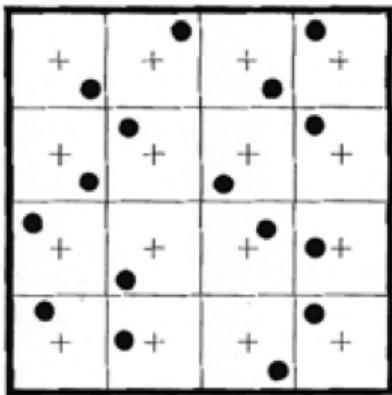
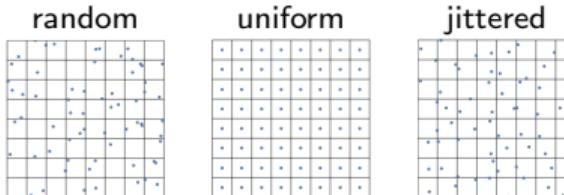
- ▶ Cone tracing. A cone (or pyramid) with apex at the eye traced through each pixel. Cone-object intersections for anti-aliasing [Amanatides 1984].
- ▶ Beam tracing. One beam for the view frustum. Split into a new beam for every polygon intersection [Heckbert and Hanrahan 1984].
- ▶ Stochastic sampling. Multiple randomly directed rays through each pixel [Cook 1986]. What is the best sampling procedure?

References

- Heckbert, P. S., and Hanrahan, P. Beam tracing polygonal objects. *Computer Graphics (SIGGRAPH 84)* 18(3), pp. 119–127. July 1984.
- Amanatides, J. Ray tracing with cones. *Computer Graphics (SIGGRAPH 84)* 18(3), pp. 129–135. July 1984.
- Cook, R. L. Stochastic sampling in computer graphics. *ACM Transactions on Graphics* 5(1), pp. 51–72. January 1986.
- Glassner, A. S. *Principles of Digital Image Synthesis*. Morgan Kaufmann. 1995.

Stratified jitter sampling

- ▶ Random numbers have a tendency to cluster.
- ▶ Stratified sampling: place samples in a grid of subpixels (substrata).
- ▶ Jitter the sample in each subpixel to avoid aliasing.



- ▶ How to make it practical when the renderer traces rays through pixel centers?
- ▶ 2D displacement vectors in the image plane. One for each subpixel.
- ▶ Using the same pattern for each pixel reduces noise.

Exercises

- ▶ Define the plane by an orthonormal basis $(\vec{b}_1, \vec{b}_2, \vec{n})$:

```
struct Onb {
    tangent: vec3f,
    binormal: vec3f,
    normal: vec3f,
};
const plane_onb = Onb(vec3f(-1.0, 0.0, 0.0), vec3f(0.0, 0.0, 1.0), vec3f(0.0, 1.0, 0.0));
```

- ▶ Texture the ground plane using an inverse mapping to obtain (u, v) -coordinates.

$$\begin{aligned} \mathbf{x} &= \mathbf{x}_0 + u\vec{b}_1 + v\vec{b}_2 \Rightarrow u = (\mathbf{x} - \mathbf{x}_0) \cdot \vec{b}_1 \\ \mathbf{x} &= \mathbf{x}_0 + u\vec{b}_1 + v\vec{b}_2 \Rightarrow v = (\mathbf{x} - \mathbf{x}_0) \cdot \vec{b}_2. \end{aligned}$$

- ▶ Perform texture sampling using nearest-neighbor and bilinear filtering.
- ▶ Use stratified jitter sampling for anti-aliasing.

Figure sources

- ▶ Woo, M., Neider, J., and Davis, T. *OpenGL Programming Guide. Second Edition. The Official Guide to Learning OpenGL, Version 1.1.* Addison-Wesley, 1997.
<http://www.openglprogramming.com/red/index.html>
- ▶ Wolfe, R. SIGGRAPH 97 Education Slide Set: Mapping Techniques. *Computer Graphics* 31(4), pp. 66-69, November 1997.
<https://education.siggraph.org/archive/slide-sets>
- ▶ Watt, A. *3D Computer Graphics*, third edition. Pearson / Addison-Wesley, 2000.
- ▶ Shreiner, D., Sellers, G., Kessenich, J., and Licea-Kane, B. *OpenGL Programming Guide. Eighth Edition. The Official Guide to Learning OpenGL, Version 4.3.* Addison-Wesley, 2013.
- ▶ RenderMan 20 Documentation. *Holdout Workflow.*
https://renderman.pixar.com/resources/RenderMan_20/risHoldOut.html
- ▶ Glassner, A. S. *Principles of Digital Image Synthesis.* Morgan Kaufmann. 1995.