

02562 Rendering - Introduction

Progressive Unidirectional Path Tracing

Jeppe Revall Frisvad


October 2023

Global illumination



©Disney Animation Studios

- ▶ Global illumination includes all light paths: $L(S|D)*E$.



Heckbert's light transport notation

- L — Light
- E — Eye
- D — Diffuse surface
- S — Specular surface
- $*$ — 0 or more interactions
- $+$ — 1 or more interactions
- $?$ — 0 or 1 interaction
- $|$ — either the path on the left or the right side

References:

- Heckbert, P. S. Adaptive radiosity textures for bidirectional ray tracing. *Computer Graphics (SIGGRAPH 90)* 24(4), pp. 145–154. 1990.

The rendering equation

- ▶ The equation we are solving is [Nicodemus 1965, Kajiya 1986]

$$L_o(\mathbf{x}, \vec{\omega}_o) = L_e(\mathbf{x}, \vec{\omega}_o) + \int_{\Omega} f_r(\mathbf{x}, \vec{\omega}_i, \vec{\omega}_o) L_i(\mathbf{x}, \vec{\omega}_i) \cos \theta_i d\omega_i,$$

- ▶ where

- L_o is observed radiance,
- L_e is emitted radiance,
- L_i is incident radiance,
- \mathbf{x} is a surface position,
- $\vec{\omega}_o$ is the direction toward the observer (direction of observation),
- $\vec{\omega}_i$ is the direction toward the light source (direction of incidence),
- f_r is the bidirectional reflectance distribution function (BRDF),
- $d\omega_i$ is a differential element of solid angle,
- Ω is the 2π solid angle around the surface normal \vec{n} at \mathbf{x} ,
- θ_i is the angle between $\vec{\omega}_i$ and the surface normal \vec{n} at \mathbf{x} , such that $\cos \theta_i = \vec{\omega}_i \cdot \vec{n}$.

References

- Nicodemus, F. E. Directional reflectance and emissivity of an opaque surface. *Applied Optics* 4(7), pp. 767–775. July 1965.
- Kajiya, J. The rendering equation. *Computer Graphics (SIGGRAPH 86)* 20(4), pp. 143–150. August 1986.

Monte Carlo integration (the sampling approach)

- ▶ The rendering equation:

$$L_o(\mathbf{x}, \vec{\omega}_o) = L_e(\mathbf{x}, \vec{\omega}_o) + \int_{2\pi} f_r(\mathbf{x}, \vec{\omega}_i, \vec{\omega}_o) L_i(\mathbf{x}, \vec{\omega}_i) \cos \theta_i \, d\omega_i .$$

- ▶ The Monte Carlo estimator:

$$L_N(\mathbf{x}, \vec{\omega}_o) = L_e(\mathbf{x}, \vec{\omega}_o) + \frac{1}{N} \sum_{j=1}^N \frac{f_r(\mathbf{x}, \vec{\omega}_{ij}, \vec{\omega}_o) L_i(\mathbf{x}, \vec{\omega}_{ij}) \cos \theta_i}{\text{pdf}(\vec{\omega}_{ij})}$$

with $\cos \theta_i = \vec{\omega}_{ij} \cdot \vec{n}$, where \vec{n} is the surface normal at \mathbf{x} .

- ▶ The Lambertian BRDF:

$$f_r(\mathbf{x}, \vec{\omega}_i, \vec{\omega}_o) = \rho_d / \pi .$$

- ▶ A good choice of pdf would be:

$$\text{pdf}(\vec{\omega}_{ij}) = \cos \theta_i / \pi .$$

Splitting the evaluation

- ▶ Distinguishing between:
 - ▶ Direct illumination L_{direct} .
 - ▶ Light reaching a surface directly from the source.
 - ▶ Indirect illumination L_{indirect} .
 - ▶ Light reaching a surface after at least one bounce.

- ▶ The rendering equation is then

$$L_o = L_e + L_{\text{direct}} + L_{\text{indirect}}.$$

- ▶ L_e is emission.
 - ▶ L_{direct} is sampling of lights.
 - ▶ L_{indirect} is sampling of the BRDF excluding lights.
- ▶ You need an emit flag in your HitInfo struct to avoid adding emission when tracing indirect illumination.

Path tracing diffuse objects

- ▶ The diffuse BRDF: $f_r = \rho_d / \pi$.
- ▶ Computing direct illumination:
 - ▶ Sample positions in random light source triangles (Δ out of n): $\text{pdf}(\mathbf{x}_j) = \frac{1}{n} \frac{1}{A_{\Delta,j}}$.
 - ▶ Estimator for L_{direct} is


$$L_{\text{direct},N} = \frac{\rho_d(\mathbf{x})}{\pi} \frac{1}{N} \sum_{j=1}^N L_e(\mathbf{x}_j \rightarrow \mathbf{x}) V(\mathbf{x}_j \leftrightarrow \mathbf{x}) \frac{\cos \theta_i \cos \theta_{\text{light}}}{\|\mathbf{x} - \mathbf{x}_j\|^2} n A_{\Delta,j}.$$

- ▶ Computing indirect illumination:
 - ▶ Set $L_e = 0$.
 - ▶ Sample directions using cosine-weighted hemisphere: $\text{pdf}(\vec{\omega}_{ij}) = \cos \theta_i / \pi$.

$$\vec{\omega}_{ij} \text{ from } (\theta_i, \phi_i) = (\cos^{-1} \sqrt{\xi_1}, 2\pi \xi_2), \quad \xi_1, \xi_2 \in [0, 1].$$

- ▶ Estimator for L_{indirect} is

$$L_{\text{indirect},N} = \rho_d(\mathbf{x}) \frac{1}{N} \sum_{j=1}^N L_i(\mathbf{x}, \vec{\omega}_{ij}).$$




Sampling a cosine-weighted hemisphere

- ▶ We sample $\perp \vec{\omega}_{ij}$ on a cosine-weighted hemisphere using spherical coordinates (θ, ϕ) in the local tangent space of surface point with normal $\vec{n} = (n_x, n_y, n_z)$.

$$(\theta, \phi) = (\cos^{-1} \sqrt{1 - \xi_1}, 2\pi \xi_2) \quad , \quad \xi_1, \xi_2 \in [0, 1) \quad (\text{random numbers}).$$

- ▶ In Euclidian coordinates, the corresponding vector in tangent space is

$$\perp \vec{\omega}_{ij} = (x, y, z) = (\cos \phi \sin \theta, \sin \phi \sin \theta, \cos \theta).$$



- ▶ We can now make a change of basis from tangent space to world space using a formula derived from quaternion rotation [Frisvad 2012, Duff et al. 2017]:

$$\vec{\omega}_{ij} = \left(x \begin{bmatrix} 1 - n_x^2 / (1 + |n_z|) \\ -n_x n_y / (1 + |n_z|) \\ -n_x \operatorname{sgn}(n_z) \end{bmatrix} + y \begin{bmatrix} -n_x n_y / (1 + |n_z|) \operatorname{sgn}(n_z) \\ (1 - n_y^2 / (1 + |n_z|)) \operatorname{sgn}(n_z) \\ -n_y \end{bmatrix} + z \begin{bmatrix} n_x \\ n_y \\ n_z \end{bmatrix} \right).$$

- ▶ Use a pseudo-random number generator to get $\xi_1, \xi_2 \in [0, 1)$ and `rotate_to_normal` for the quaternion rotation.

What is a quaternion?

- ▶ Generalization of complex numbers for representation of rotations.
- ▶ Specified as a 4-vector with a vector imaginary part \mathbf{q}_v and a scalar real part q_w :

$$\hat{\mathbf{q}} = (\mathbf{q}_v, q_w) = iq_x + jq_y + kq_z + q_w, \quad \mathbf{q}_v = (q_x, q_y, q_z), \quad i^2 = j^2 = k^2 = ijk = -1.$$

Multiplication: $\hat{\mathbf{q}}\hat{\mathbf{r}} = (\mathbf{q}_v \times \mathbf{r}_v + r_w \mathbf{q}_v + q_w \mathbf{r}_v, q_w r_w - \mathbf{q}_v \cdot \mathbf{r}_v).$


Addition: $\hat{\mathbf{q}} + \hat{\mathbf{r}} = (\mathbf{q}_v + \mathbf{r}_v, q_w + r_w).$

Conjugate: $\hat{\mathbf{q}}^* = (-\mathbf{q}_v, q_w).$

Norm: $\text{norm}(\hat{\mathbf{q}}) = \sqrt{\hat{\mathbf{q}}\hat{\mathbf{q}}^*} = \sqrt{\mathbf{q}_v \cdot \mathbf{q}_v + q_w^2}.$

Identity: $\hat{\mathbf{i}} = (0, 1).$

Inverse: $\hat{\mathbf{q}}^{-1} = \hat{\mathbf{q}}^* / [\text{norm}(\hat{\mathbf{q}})]^2.$




- ▶ Unit quaternion specifying rotation of 2ϕ radians around an axis \mathbf{u}_q :

$$\hat{\mathbf{q}} = (\sin\phi \mathbf{u}_q, \cos\phi).$$

- ▶ Use homogeneous coordinates to convert a position or direction vector $\mathbf{p} = (p_x, p_y, p_z, p_w)$ to a quaternion $\hat{\mathbf{p}}$.
- ▶ Apply quaternion rotation using $\hat{\mathbf{q}}\hat{\mathbf{p}}\hat{\mathbf{q}}^{-1}$. For a unit quaternion, use $\hat{\mathbf{q}}\hat{\mathbf{p}}\hat{\mathbf{q}}^*.$

Rotating from tangent space to world space




- ▶ Unit quaternion specifying rotation from \mathbf{u} to \mathbf{v} :

$$\hat{\mathbf{q}} = \left(\sin \frac{\theta}{2} \mathbf{w}, \cos \frac{\theta}{2} \right) = \left(\frac{\mathbf{u} \times \mathbf{v}}{\sqrt{2(1 + \mathbf{u} \cdot \mathbf{v})}}, \frac{1}{2} \sqrt{2(1 + \mathbf{u} \cdot \mathbf{v})} \right).$$

- ▶ Let us set $\mathbf{u} = (0, 0, 1)$ and $\mathbf{v} = \vec{n} = (n_x, n_y, n_z)$, then

$$\hat{\mathbf{q}} = \left(\frac{(-n_y, n_z, 0)}{\sqrt{2(1 + n_z)}}, \frac{1}{2} \sqrt{2(1 + n_z)} \right).$$



- ▶ Applying $\hat{\mathbf{q}}$ to a direction vector $\perp \vec{\omega} = (\omega_x, \omega_y, \omega_z)$, we rotate from tangent to world space: $\hat{\mathbf{q}} (\perp \vec{\omega}, 0) \hat{\mathbf{q}}^*$.
- ▶ This simplifies to the formula for rotate_to_normal [Frisvad 2012, Duff et al. 2017].

References:

- Frisvad, J. R. Building an orthonormal basis from a 3D unit vector without normalization. *Journal of Graphics Tools* 16(3), pp. 151–159. August 2012.
- Duff, T., Burgess, J., Christensen, P., Hery, C., Kensler, A., Liani, M., and Villemin, R. Building an orthonormal basis, revisited. *Journal of Computer Graphics Techniques* 6(1), pp. 1–8. March 2017.

Progressive unidirectional path tracing

1. **Generate rays** from the eye through pixel positions
 2. **Trace the rays** and evaluate the rendering equation for each ray.
 3. **Randomize the position** within the pixel area to Monte Carlo integrate (measure) the radiance arriving in a pixel.
- ▶ Path tracing is the idea of using $N = 1$ in the Monte Carlo estimators for the reflected radiance to generate a path instead of a tree.
 - ▶ Noise is reduced by progressive updates of the measurement.
 - ▶ Update the rendering result in a pixel L_j after rendering a new frame with result L_{new} using

$$L_{j+1} = \frac{L_{\text{new}} + jL_j}{j+1} .$$

- ▶ Progressive (stop and go) rendering is convenient for several reasons:
 - ▶ No need to start over.
 - ▶ Result can be stored and refined later if need be.
 - ▶ Convergence can be inspected during progressive updates.

Stopping paths probabilistically (Russian roulette)

- ▶ Do either one thing or another (not both).
- ▶ Sample an event.
 - ▶ Example: Reflection or absorption.
 - ▶ Example: Reflection or transmission.
- ▶ How? Sample a step function.
- ▶ What are the steps? The probabilities that events occur.
 - ▶ Example: Either diffuse reflection or absorption.
 - ▶ What are the probabilities?
 - ▶ The simplest idea is the average of the diffuse reflectance:

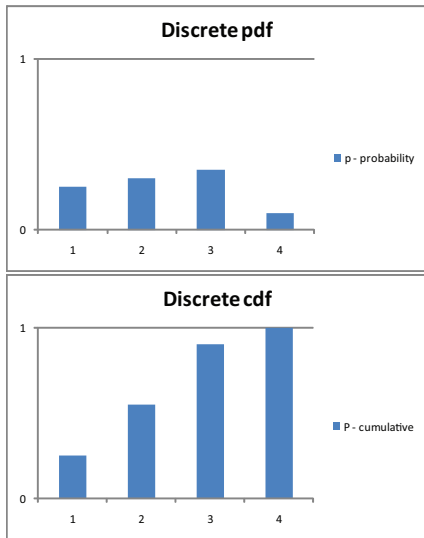
$$\text{probability of diffuse reflection} = \frac{\rho_{d,R} + \rho_{d,G} + \rho_{d,B}}{3} .$$

- ▶ The probability of absorption is then one minus the probability of diffuse reflection (the else-case).
- ▶ In the case of diffuse reflection, a ray with sampled direction is traced to evaluate the indirect illumination.


Sampling a step function (Russian roulette)

Algorithm:


```
sample  $\xi \in [0, 1]$  uniformly;  
if ( $\xi < P_1$ )  
    call event 1;  
    divide by  $p_1$ ;  
else if ( $\xi < P_2$ )  
    call event 2;  
    divide by  $p_2$ ;  
else if ( $\xi < P_3$ )  
    ...  
else if ( $\xi < P_4$ )  
    ...
```




Direct versus global illumination (and jittered versus progressive)



progressive
1 minute
400 samples



jittered
1 minute
484 samples



progressive
13.4 minutes
1000 samples

- The worst case convergence is an error proportional to $\frac{1}{\sqrt{N}}$, where N is the number of samples.

How to get pseudo-random numbers on the GPU?

- ▶ We need a seeded pseudo-random number generator (PRNG) and a different seed for each thread.
- ▶ In rendering, a natural choice is one thread for each pixel in each frame.
- ▶ Given pixel index and frame number, we can use Marsaglia's xorshift bit scrambling to get a pseudo-random seed for the PRNG:

```
// PRNG xorshift seed generator by NVIDIA
fn tea(val0: u32, val1: u32) -> u32
{
  const N = 16u; // User specified number of iterations
  var v0 = val0; var v1 = val1; var s0 = 0u;
  for(var n = 0u; n < N; n++) {
    s0 += 0x9e3779b9;
    v0 += ((v1<<4)+0xa341316c)^(v1+s0)^((v1>>5)+0xc8013ea4);
    v1 += ((v0<<4)+0xad90777d)^(v0+s0)^((v0>>5)+0x7e95761e);
  }
  return v0;
}
```

- ▶ With a seed, we can use a multiplicative congruential PRNG:

$$Z_i = AZ_{i-1} \quad \text{mod } M$$

if we can find a good multiplier A for a modulus $M \in [0, 2^{32})$.

The good multiplier for a multiplicative congruential generator

- ▶ The good multiplier should have a full period (reach all integers from 0 to $M - 1$) and exhibit good equidistribution if we use it to sample points in a k D hypercube.
- ▶ Fishman [1990] performed an exhaustive analysis to find the best multiplier for $M = 2^{32}$. This is good if we want random numbers $\xi \in [0, 1]$.
- ▶ However, we want $\xi \in [0, 1)$ (e.g. when sampling an index $i < n$ by $i = \lfloor \xi n \rfloor$).
- ▶ Hui-Chin Tang [2007] performed an exhaustive search to find the best multiplier for $M = 2^{31} - 1$ with good equidistribution (for $1 < k \leq 8$). Using $M = 2^{31}$:

```
// Generate random unsigned int in [0, 2^31)
fn mcg31(prev: ptr<function, u32>) -> u32
{
    const LCG_A = 1977654935u; // Multiplier from Hui-Ching Tang [EJOR 2007]
    *prev = (LCG_A * (*prev)) & 0x7FFFFFFF;
    return *prev;
}

// Generate random float in [0, 1)
fn rnd(prev: ptr<function, u32>) -> f32
{
    return f32(mcg31(prev)) / f32(0x80000000);
}
```

Sampling a triangle mesh


- ▶ Uniformly sample a triangle index $i = \lfloor \xi n \rfloor$ ($\text{pdf}(\Delta) = 1/n$, where n is the number of triangles in the mesh).
- ▶ Uniformly sample a position on the triangle ($\text{pdf}(\mathbf{x}_{\ell,j,\Delta}) = 1/A_\Delta$, where A_Δ is the triangle area):

1. Sample barycentric coordinates $(\alpha, \beta, \gamma = 1 - \alpha - \beta)$:

$$\alpha = 1 - \sqrt{\xi_1}$$

$$\beta = (1 - \xi_2)\sqrt{\xi_1}$$

$$\gamma = \xi_2\sqrt{\xi_1}$$



2. Use the barycentric coordinates for linear interpolation of triangle vertices $(\mathbf{q}_0, \mathbf{q}_1, \mathbf{q}_2)$ to obtain a point in the triangle: $\mathbf{x}_{\ell,j} = \alpha \mathbf{q}_0 + \beta \mathbf{q}_1 + \gamma \mathbf{q}_2$.
 3. Use the barycentric coordinates for linear interpolation of triangle vertex normals to obtain the normal $\vec{n}_{\ell,j}$ at the sampled surface point. Normalize after interpolation.
- ▶ The pdf of this sampling of the surface is: $\text{pdf}(\mathbf{x}_{\ell,j}) = \text{pdf}(\Delta)\text{pdf}(\mathbf{x}_{\ell,j,\Delta}) = \frac{1}{n} \frac{1}{A_\Delta}$.

Soft shadows

- Sampler:


$$\vec{\omega}'_j = \frac{\mathbf{x}_{\ell,j} - \mathbf{x}}{\|\mathbf{x}_{\ell,j} - \mathbf{x}\|}.$$

- From solid angle to area:

$$d\omega' = \frac{\cos \theta_\ell}{r^2} dA = \frac{\vec{n}_{\ell,j} \cdot (-\vec{\omega}'_j)}{\|\mathbf{x}_{\ell,j} - \mathbf{x}\|^2} dA.$$

- Using the Lambertian BRDF, $f_r = \rho_d/\pi$, and triangle mesh sampling of a point $\mathbf{x}_{\ell,j}$ on the light, $\text{pdf}(\mathbf{x}_{\ell,j}) = 1/(nA_\Delta)$, the Monte Carlo estimator for area lights is:

$$\begin{aligned} L_N(\mathbf{x}, \vec{\omega}) &= \frac{1}{N} \sum_{j=1}^N \frac{f_r(\mathbf{x}, \vec{\omega}'_j, \vec{\omega}) L_i(\mathbf{x}, \vec{\omega}'_j) \cos \theta \cos \theta_\ell}{\text{pdf}(\mathbf{x}_{\ell,j}) r^2} \\ &= \frac{\rho_d(\mathbf{x})}{\pi} \frac{1}{N} \sum_{j=1}^N \underbrace{L_e(\mathbf{x}_{\ell,j}, -\vec{\omega}'_j) V(\mathbf{x}, \mathbf{x}_{\ell,j}) \frac{\vec{n}_{\ell,j} \cdot (-\vec{\omega}'_j)}{\|\mathbf{x}_{\ell,j} - \mathbf{x}\|^2} nA_\Delta (\vec{\omega}'_j \cdot \vec{n})}_{\text{returned as } L \text{ from area light sampler}}. \end{aligned}$$



Exercises

- ▶ Use ping-pong rendering for progressive updates with sampling of a random positions in each pixel (replaces stratified jitter sampling).
- ▶ Implement area light triangle mesh sampling to render soft shadows.
- ▶ Implement sampling of a cosine-weighted hemisphere.
- ▶ Write a shader for Lambertian materials that includes path-traced indirect illumination.

Render-to-texture and ping-ponging (pipeline)

- The pipeline changes: we add an f32 texture target for our render results.

```
const pipeline = device.createRenderPipeline({  
  layout: "auto",  
  vertex: {  
    module: wgs1,  
    entryPoint: "main_vs",  
  },  
  fragment: {  
    module: wgs1,  
    entryPoint: "main_fs",  
    targets: [ { format: canvasFormat },  
               { format: "rgba32float" } ]  
  },  
  primitive: {  
    topology: "triangle-strip",  
  },  
});
```

Render-to-texture and ping-ponging (textures)

- ▶ Create one texture as a render target and one to load the previous result from.

```
let textures = new Object();
textures.width = canvas.width;
textures.height = canvas.height;
textures.renderSrc = device.createTexture({
  size: [canvas.width, canvas.height],
  usage: GPUTextureUsage.RENDER_ATTACHMENT | GPUTextureUsage.COPY_SRC,
  format: 'rgba32float',
});
textures.renderDst = device.createTexture({
  size: [canvas.width, canvas.height],
  usage: GPUTextureUsage.TEXTURE_BINDING | GPUTextureUsage.COPY_DST,
  format: 'rgba32float',
});
```

- ▶ The render target is a source for copying data to the destination texture binding.
- ▶ Use no sampler for the texture binding.

Render-to-texture and ping-ponging (render pass)

- The render pass changes: we include a second color attachment and a texture-to-texture copy.

```
function render(device, context, pipeline, textures, bindGroup)
{
  const encoder = device.createCommandEncoder();
  const pass = encoder.beginRenderPass({
    colorAttachments: [
      { view: context.getCurrentTexture().createView(), loadOp: "clear", storeOp: "store" },
      { view: textures.renderSrc.createView(), loadOp: "load", storeOp: "store" } ]
  });
  pass.setPipeline(pipeline);
  pass.setBindGroup(0, bindGroup);
  pass.draw(4);
  pass.end();

  encoder.copyTextureToTexture({ texture: textures.renderSrc }, { texture: textures.renderDst },
    [textures.width, textures.height]);

  // Finish the command buffer and immediately submit it.
  device.queue.submit([encoder.finish()]);
}
```

Render-to-texture and ping-ponging (fragment shader)

- ▶ The fragment shader implements the progressive updating.

```
struct FSOut {  
    @location(0) frame: vec4f,  
    @location(1) accum: vec4f,  
};  
  
@fragment  
fn main_fs(@builtin(position) fragcoord: vec4f, @location(0) coords: vec2f) -> FSOut  
{  
    let launch_idx = u32(fragcoord.y)*uniforms_ui.width + u32(fragcoord.x);  
    var t = tea(launch_idx, uniforms_ui.frame);  
    let jitter = vec2f(rnd(&t), rnd(&t))/f32(uniforms_ui.height);  
    :  
  
    // Progressive update of image  
    let curr_sum = textureLoad(renderTexture, vec2u(fragcoord.xy), 0).rgb*f32(uniforms_ui.frame);  
    let accum_color = (result + curr_sum)/f32(uniforms_ui.frame + 1u);  
  
    var fsOut: FSOut;  
    fsOut.frame = vec4f(pow(accum_color, vec3f(1.0/uniforms_f.gamma)), 1.0);  
    fsOut.accum = vec4f(accum_color, 1.0);  
    return fsOut;  
}
```

- ▶ The texture binding is in renderTexture. The texture target is in location 1 of the output. The frame resolution and number are uploaded as uniforms.